

JavaLife

William Rossi

CS 1699 - DELIVERABLE 4: Performance Testing Conway's Game of Life

Profiling Conway's Game of Life turned out to be a much more problematic task than I had originally thought it would be. I performed a good deal of troubleshooting to figure out how to run the program with the default packages included (curse you, com.laboony!) and then it took even further tinkering before I figured out how to get VisualVM to profile my CPU at the method level. What I learned is that I had to tell the profiler what package to look in for the methods that it would be profiling (CURSE YOU, COM.LABOON!).

Once I finished the initial setup, however, the profiling process was actually remarkably easy. I set the game to run through 30,000 iterations and profiled the CPU usage of the game's methods while it was running, taking a snapshot after an arbitrary (but long enough) amount of time had passed. The snapshots below show what methods were being called, how many times they were called, and how much time the CPU spent executing these methods (relative to the other methods). The results were pretty striking, and it was immediately apparent to me that the `toString()` method in the `World` class was the most inefficient method. As shown in the snapshot before I refactored the method, the CPU was spending just about all of its time (100%, although I'm fairly certain that's not an exact figure) executing the `toString()` method, even though other methods--like the `getNumNeighbors()` method--were being invoked much more often.

Upon examination of the `World` class, I realized that its inefficiency stemmed from its use of string concatenation: in order to create the visual aspect of the game, the `toString()` method was creating an initial string and then continuously concatenating it with new strings. From my prior experience with the Java language I knew that string objects were immutable, so what this method was really doing was creating a new string object every time the `"+="` operator was used to concatenate the old string with a new one. This concatenation happened many times within each invocation of the `toString()` method, and the `toString()` method was being called (in total) 30,000 times! Not only did this constant creation of new objects waste time, but it also wasted space!

To refactor the method, I decided to replace the string object within the method with a `StringBuilder`. Then, instead of concatenating two strings together repeatedly, I could simply append the necessary strings onto the `StringBuilder`. Finally, I would return the `String` representation of that `StringBuilder` (by calling its built-in `toString()` method) so that the type of variable returned would not need to be changed--and this would also ensure that my unit tests were passing both before and after the refactor as is also shown in the screenshots below. This not only saved space because each call to `World.toString()` would only create ONE `StringBuilder` object rather than MANY `String` objects of increasing size, but it also saved time because appending to a `StringBuilder` is much faster than creating a new `String` by concatenating two strings. As shown in the "after" snapshot of the CPU profile, the `toString()` method decreased in

total execution time by about 80%, and note that the method itself had been invoked MANY more times than in the snapshot taken before the refactor!

To unit test the method, I tried to think of various boundary cases that might apply to the method. My first thought was to test the `toString()` output of a world of size zero, and this test was easy to write because I could very quickly determine what the output should be and how long it should be. I wanted to test a very large world, but manually determining the size and value of that `toString()` output was too challenging. I worked around this challenge by testing that two identical worlds generated from the same arbitrary arguments both produced the same output and output length when their `toString()` methods were invoked; the worlds ranged in size from 0 to 199. Once that was done, I wanted to ensure that two worlds produced using different sizes would NOT produce strings of the same length when their `toString()` methods were invoked (as this would mean that the program was broken), so I tested that as well. Finally, I decided to test that a world created from arbitrary (but correct in terms of the program's usage) values would not return a null string from its `toString()` method, because this would also mean that the program was broken. Since the pre- and post-refactor `toString()` method both returned `String` objects (and these strings were functionally equivalent), the tests passed before and after I refactored the method.

Aside from the challenges I faced initially (described in the first paragraph), there was one minor annoyance that I faced. I realized that, even though I knew from experience how `String` concatenation worked in terms of space usage, it would be more convincing if I profiled the memory usage of the program both before and after my refactor to demonstrate that I had, in fact, decreased the amount of space used by the program. Unfortunately, for some reason I could not make my machine run the original version of the `World.toString()` method; I think maybe it had cached the new version and continued to run that. I tried clearing my cache, and still the same problem was occurring. For whatever reason, every time I tried to memory profile the old version of the program, it kept creating `StringBuilder` objects, and this was a tipoff to me that it was not calling the old `toString()` method because that did not contain any `StringBuilders`. I eventually became so frustrated with the problem that I gave up. I figured it would be pointless to put in a profile of the memory usage of the newly refactored method since there was nothing to compare it to, AND I am also confident that everything I've already mentioned about the space savings of `StringBuffers` versus `Strings` when concatenating is true.

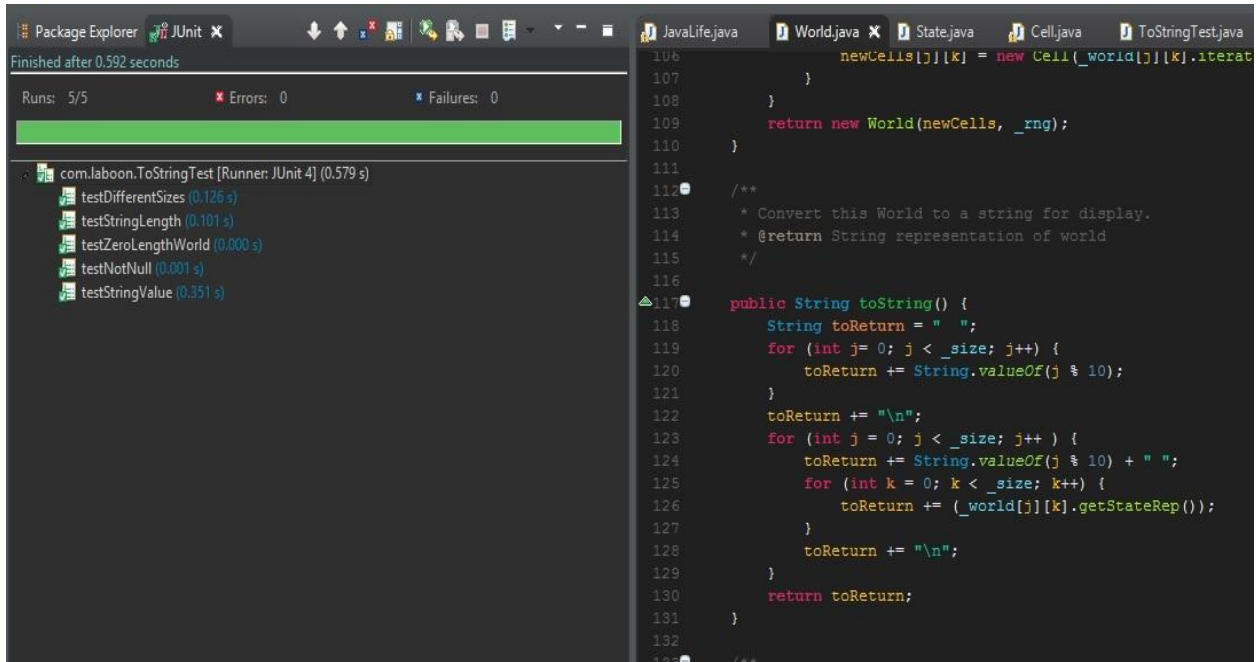
Before refactor:

Call Tree - Method	Time [%] ▼	Time	Invocations
main		25540 ms (100%)	1
└─ com.laboan.World.toString ()		26201 ms (100%)	8275
└─ Self time		26272 ms (100%)	8275
└─ com.laboan.Cell.getStateRep ()		0.000 ms (0%)	20686185
└─ com.laboan.Cell.getStateRep ()		0.094 ms (0%)	100
└─ com.laboan.World.iterate ()		0.000 ms (0%)	8275
└─ Self time		510 ms (2%)	8275
└─ com.laboan.World.<init> (com.laboan.Cell)		1.12 ms (0%)	8275
└─ com.laboan.World.getNumNeighbors (com.laboan.Cell)		0.000 ms (0%)	20687500
└─ Self time		1001 ms (3.9%)	20687500
└─ com.laboan.Cell.isAlive ()		0.000 ms (0%)	165500000
└─ com.laboan.Cell.iterate (int)		0.000 ms (0%)	20687500
└─ com.laboan.Cell.<init> (com.laboan.State)		0.000 ms (0%)	20687500

After refactor:

Call Tree - Method	Time [%] ▼	Time	Invocations
main		3806 ms (100%)	1
└─ com.laboan.World.iterate ()		3022 ms (79.4%)	17491
└─ Self time		1725 ms (45.3%)	17491
└─ com.laboan.World.getNumNeighbors (com.laboan.Cell)		1460 ms (38.4%)	43727009
└─ Self time		4076 ms (100%)	43727009
└─ com.laboan.Cell.isAlive ()		0.000 ms (0%)	349816067
└─ com.laboan.World.<init> (com.laboan.Cell)		3.78 ms (0.1%)	17490
└─ com.laboan.Cell.iterate (int)		0.000 ms (0%)	43727008
└─ com.laboan.Cell.<init> (com.laboan.State)		0.000 ms (0%)	43727008
└─ com.laboan.World.toString ()		785 ms (20.6%)	17490
└─ Self time		931 ms (24.5%)	17490
└─ com.laboan.Cell.getStateRep ()		0.000 ms (0%)	43725000

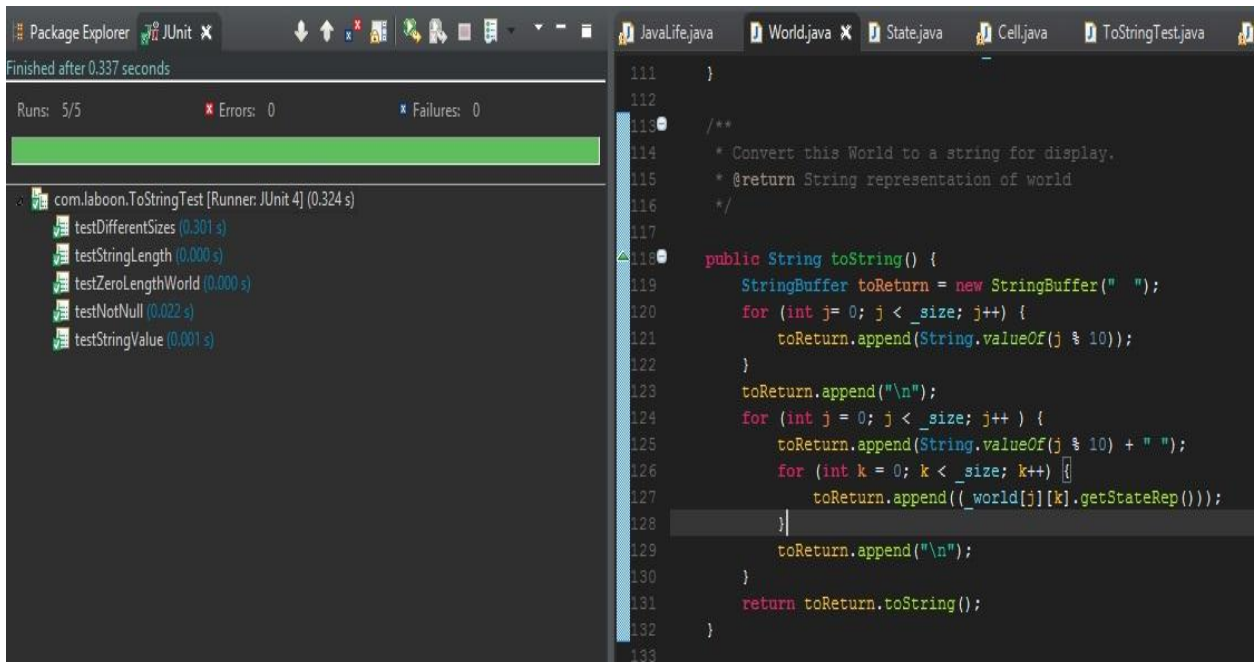
Unit Tests before refactor:



The screenshot shows an IDE with the JUnit test runner on the left and the source code on the right. The test runner indicates that 5 tests passed, with 0 errors and 0 failures, taking 0.579 seconds. The tests listed are: testDifferentSizes (0.126 s), testStringLength (0.101 s), testZeroLengthWorld (0.000 s), testNotNull (0.001 s), and testStringValue (0.351 s). The source code on the right shows the `World` class with a `toString` method that uses multiple nested loops and string concatenation to build the world's string representation.

```
106         newCells[j][k] = new Cell(_world[j][k].iterat
107     }
108 }
109 return new World(newCells, _rng);
110 }
111
112 /**
113  * Convert this World to a string for display.
114  * @return String representation of world
115  */
116
117 public String toString() {
118     String toReturn = " ";
119     for (int j= 0; j < _size; j++) {
120         toReturn += String.valueOf(j % 10);
121     }
122     toReturn += "\n";
123     for (int j = 0; j < _size; j++) {
124         toReturn += String.valueOf(j % 10) + " ";
125         for (int k = 0; k < _size; k++) {
126             toReturn += (_world[j][k].getStateRep());
127         }
128         toReturn += "\n";
129     }
130     return toReturn;
131 }
132
133 /**
```

Unit Tests after refactor:



The screenshot shows the same IDE after refactoring. The test runner now shows the tests completed in 0.324 seconds, with the same 5 tests passing. The `toString` method in the `World` class has been refactored to use a `StringBuffer` for building the string, which is more efficient than concatenation. The code is cleaner and easier to read.

```
111     }
112
113     /**
114     * Convert this World to a string for display.
115     * @return String representation of world
116     */
117
118     public String toString() {
119         StringBuffer toReturn = new StringBuffer(" ");
120         for (int j= 0; j < _size; j++) {
121             toReturn.append(String.valueOf(j % 10));
122         }
123         toReturn.append("\n");
124         for (int j = 0; j < _size; j++) {
125             toReturn.append(String.valueOf(j % 10) + " ");
126             for (int k = 0; k < _size; k++) {
127                 toReturn.append((_world[j][k].getStateRep()));
128             }
129             toReturn.append("\n");
130         }
131         return toReturn.toString();
132     }
133 }
```